

Kyffin Williams: Digital Analysis of Paintings

Final Report for CS39440 Major Project

Author: Alexander David Brown (adb9@aber.ac.uk)

Supervisor: Hannah Dee (hmd1@aber.ac.uk)

22nd April 2012

Version: 0.0.1067 (Draft)

This report was submitted as partial fulfilment of a MEng degree in
Software Engineering (G601)

Department of Computer Science
Aberystwyth University
Aberystwyth
Ceredigion
SY23 3DB
Wales, UK

Declaration of originality

In signing below, I confirm that:

- This submission is my own work, except where clearly indicated.
- I understand that there are severe penalties for plagiarism and other unfair practice, which can lead to loss of marks or even the withholding of a degree.
- I have read the sections on unfair practice in the Students' Examinations Handbook and the relevant sections of the current Student Handbook of the Department of Computer Science.
- I understand and agree to abide by the University's regulations governing these issues.

Signature

Date

Consent to share this work

In signing below, I hereby agree to this dissertation being made available to other students and academic staff of the Aberystwyth Computer Science Department.

Signature

Date

Acknowledgements

Abstract

CONTENTS

1	Background & Objectives	1
1.1	Sir John “Kyffin” Williams	1
1.2	Interdisciplinary work with the National Library of Wales	1
1.2.1	Continuation of the Kyffin Project	2
1.3	Existing Work	3
1.3.1	Edge-Orientated Gradients	3
1.3.2	Brush-stroke Analysis	4
1.4	Analysis Objectives	5
1.4.1	Colour-space Analysis	5
1.4.2	Texture Analysis	6
1.4.3	Brush-stroke Analysis	6
1.4.4	Ensemble Techniques	6
1.5	Classification Objectives	7
1.5.1	Classification	7
1.5.2	Exemplars	8
2	Development Process	9
2.1	Introduction	9
2.2	Modifications	9
3	Design	10
3.1	Methodology	10
3.2	Overall Architecture	11
4	Implementation	13
4.1	Basic Structure	13
4.1.1	Loading Data	13
4.1.2	Top-Level Classes	13
4.1.3	Command Line Interface	14
4.2	Colour Space Analysis	14
4.2.1	Colour Models	14
4.2.2	Colour Histograms	15
4.3	Texture Analysis	15
4.3.1	Edge Count	15
4.3.2	Edge Orientation	15
4.3.3	Steerable Filters	16
4.3.4	Gabor Filters	16
4.4	Brush-Stroke Analysis	17
4.5	Classification and Validation	17
4.5.1	k -Nearest Neighbour	17
4.5.2	Leave-One-Out Cross Validation	17
4.5.3	Weka 3	17
4.5.4	Exemplars	17
4.5.5	Statistically Classified Exemplars	18
4.5.6	cv2 and Histogram Results	18
4.6	3 rd Party Libraries and Tools	19

4.6.1 Python	19
4.6.2 OpenCV	19
4.6.3 scipy & numpy	19
4.6.4 matplotlib	19
4.6.5 Weka 3	19
4.6.6 git & github	19
5 Testing	20
5.1 Overall Approach to Testing	20
5.2 Validation	20
5.2.1 Leave-One-Out Cross Validation	20
5.2.2 Validation using Weka	20
6 Evaluation	21
6.1 Evaluation of Requirements	21
6.2 Evaluation of Design	21
6.3 Evaluation of Tools	21
6.3.1 Programming Language	21
6.3.2 Image Processing/Computer Vision Libraries	21
6.3.3 Machine Learning Libraries	21
6.3.4 Scientific and Numeric Libraries	21
Appendices	22
A Paper Submission for ISPA 2013	23
B 3rd Party Libraries and Tools	29
2.1 Python 2.7	29
2.1.1 setuptools	29
2.1.2 scipy	29
2.1.3 numpy	29
2.1.4 matplotlib	29
2.1.5 liac-arff	29
2.2 OpenCV	29
2.2.1 OpenCV Python	29
2.3 Weka 3	29
2.4 git	30
2.4.1 github	30
C Equations	31
3.1 Statistical Equations	31
3.1.1 Mean	31
3.1.2 Standard Deviation	31
3.1.3 Pearson's product-moment coefficient	31
3.2 Distance Equations	31
3.2.1 Manhattan Distance	31
3.2.2 Euclidean Distance	31
3.3 Filter Equations	31
3.3.1 Gradient Direction	31

3.3.2 Discrete Derivative Masks	32
3.3.3 Gabor Filter	32
D Code Samples	33
4.1 Example CSV Parsing Code	33
4.2 argparse Example Code	33
4.3 Gabor Filter Example Implementation	33
4.4 OpenCV Histogram Example Code	34
E Spreadsheet Data	36
Annotated Bibliography	38

LIST OF FIGURES

1.1	Example of a 1 by 2 Discrete Derivative Mask	3
1.2	Example of a 1 by 3 Discrete Derivative Mask	3
1.3	Pseudocode for Leave-One-Out Cross Validation	7
1.4	Pseudocode for k -Nearest Neighbour	7
3.1	Overview of the Classification Methodology	10
3.2	Basic Overall Architecture	11
3.3	Overall Architecture with Factory Methods	11
3.4	Overall architecture with Interfaces for Analysers and Classifiers	12
4.1	Painting Data in CSV Format	13
4.2	Steerable Filters	16
4.3	Nearest Exemplar Classification Psuedocode	18
4.4	Psuedocode for generating a Statistically Classified Exemplar	18

LIST OF TABLES

4.1	Layout of the Painting Data Spreadsheet	13
4.2	Command Line Arguments	14
4.3	Layout of the Exemplar Spreadsheet	17
E.1	Exemplar Spreadsheet Data	37

LIST OF LISTINGS

D.1	Example CSV Parsing Code from http://docs.python.org/2/library/csv.html	33
D.2	Example argparse Code from http://docs.python.org/2/library/argparse.html . . .	33
D.3	Example implementation of a Gabor Filter in MATLAB from wikipedia [8] . . .	34
D.4	Example Histogram calculation and displaying code from OpenCV [4].	34

Chapter 1

Background & Objectives

1.1 Sir John “Kyffin” Williams

Sir John “Kyffin” Williams (1918-2006) was a Welsh painter and printmaker, widely regarded as the defining artist of Wales during the 20th century [10]. He was advised to take up art by a doctor after failing a British Army medical examination because of an ‘abnormality’ (epilepsy) as something which would not tax his brain.

He studied at the Slade School of Fine Art and taught art in Highgate School, after which he retired to Anglesey until he died in 2006 after a long battle with cancer.

His most characteristic pictures are of Welsh landscapes, painted with thick layers of oil paint applied with a palette knife [7]. Most of his paintings are highly textural; to the point of being 3-dimensional.

As his life progressed Kyffin’s ‘abnormality’ grew steadily worse, especially when exposed to bright light. As a result most of his paintings are of overcast Welsh landscapes and tend to become visibly darker over time [13]. By eye it is generally quite easy to approximate the time period in which a painting was created.

In 1969 he won a scholarship to study and paint in Y Wladfa; the Welsh settlement in Patagonia. This period of his life is very obvious from his paintings as there is a complete contrast in colour between Patagonian and Welsh landscapes.

1.2 Interdisciplinary work with the National Library of Wales

This project was initially suggested through a conversation between Hannah Dee and Gareth “Llyod” Roderick about image processing and art. Llyod is a PhD student at the National Library of Wales (NLW). Their initial idea was to try to geolocate a Kyffin painting on a map to build up a geographical representation of Kyffin’s work.

Hannah started to create a prototype for performing geographical analysis, this proved to be a difficult task and one which is still being researched.

However, the nature of Kyffin’s illness and painting style allows for a second form of analysis: temporal. As previously stated it is fairly easy to judge by eye a good approximation of the period in which a Kyffin painting was created. It should, therefore, follow that this process can be performed digitally.

When I started this project I was given a “database” (in reality this was just a spreadsheet) Llyod had produced, containing information of Kyffin Williams’ paintings, including: title, year,

category (landscape, portrait, etc.), canvas size and a few additional details which aren't so relevant to the project.

The first meeting held was between Llyod, Hannah and I, in which we discussed the current state of the project, what our aims for the project were and what form of help Llyod could provide to us. As one of the objectives of this project is to, eventually, get a paper published, the relevant details of the process we would need to go through if we wanted to do so.

The second meeting was between Hannah, Llyod, Lorna M. Hughes (Llyod's supervisor) and I. Again we discussed the state of the project. Llyod had also produced a better version of his "database" to be more machine readable and succinct. A lot of information came from this meeting;

- The "cut-off" point between early and late is around 1973.
- The size of the canvas might be a useful data point to use in classification, as Kyffin sold more paintings he would have had the money available for larger canvases and the paint for said canvas.
- It is a little dubious as to whether some dates can be trusted. One painting owned by the NLW was stated to be his last painting, but Lorna believes it was painted much earlier and claimed to be his last to improve the sale price.
- Llyod may have found date markings on some paintings. These again may not be accurate, but may prove to increase the sample size.
- It should be easy to provide a "no later than" estimate for each painting from the art historians.
- Paul (?) should be able to produce some exemplars for us as a ground truth.
- Llyod may be able to find more paintings in the hands of private collectors to increase the sample size.
- Llyod had been playing around with ImageJ to do some basic graph plotting. This might be useful to look at further to expand my own work.

There were also more detailed discussions about publications, particularly in a digital humanities journal.

1.2.1 Continuation of the Kyffin Project

There are several projects that could continue on from the Kyffin Project.

One was to use the Learning/Teaching development fund to produce a web-based front-end for of some of my analysis.

Another venture was to look into PhD funding to build up a 3D map of some of Kyffin's paintings and being able to display it (perhaps via HTML5 and WebGL) so they can explore the painting digitally how it is meant to be in real life.

1.3 Existing Work

1.3.1 Edge-Orientated Gradients

As Kyffin Williams' work is highly texturally, looking at the edge orientation of the image is likely to be a valuable technique to use.

One technique recommended by Hannah was to look at Histogram of Orientated Gradients (HOG). The suggested paper outlined the use of grids of HOG descriptors to improve the feature set for robust visual object recognition [9]. As it significantly improves the feature set it seems sensible to try and implement it as a technique without the Kyffin project to experiment with a non colour-based approach.

The approach involves quite a few separate steps, only some of which are relevant to the project:

1. Gamma and colour normalization. Grayscale, Red, Green, Blue (RGB) and L, a, b Colour Space (LAB) spaces were used. RGB and LAB give similar results. Grayscale reduced performance less than square root gamma compression, but not as much as log compression.
2. Compute gradients. Often the simplest are the better here; Gaussian smoothing followed by discrete derivative masks (e.g.: figure 1.1, figure 1.2), etc. For colour this was done for each channel, and take the one with the largest norm.
3. Spatial and Orientation binning:
 - Spatial binning is done by splitting the images into cells which can be rectangle or radial.
 - Orientation binning are spaced equally between either 0-180 "unsigned" or 0-360 "signed" bins.
4. Normalisation and Descriptor Blocks. Gradients vary over foreground/background, etc. Typically the blocks were overlapped so that each scalar response contains several components.
5. Pass a detector window across the image.
6. Run through a Linear Support Vector Machine (SVM) to classify the image.

$$\begin{bmatrix} -1 \\ 1 \end{bmatrix}$$

Figure 1.1: Example of a 1 by 2 Discrete Derivative Mask

$$\begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}$$

Figure 1.2: Example of a 1 by 3 Discrete Derivative Mask

Obviously, when applying this as an analysis technique to paintings, there are some points which are completely irrelevant. Passing a detector window and running through a Linear SVM are the obvious two. Normalisation is another unneeded step; computing performance isn't likely to be a large issue for this project; so long as the techniques complete within a semi-reasonable amount of time.

This leaves the act of computing the gradients (again, this can be done without Gaussian smoothing as that reduces accuracy) which should be a simple matter implemented by earlier techniques. Binning by Rectangular Histogram of Orientated Gradients (R-HOG) or Circular Histogram of Orientated Gradients (C-HOG) descriptors, which may prove to be one of the more difficult parts of implementing this technique.

R-HOG descriptors have similarities to Scale-Invariant Feature Transform (SIFT) descriptors, but are used quite differently; SIFT descriptors are optimised for sparse baseline matching whilst R-HOG descriptors are optimised for the dense and robust coding of a spatial form. The size of the descriptor affects performance when using R-HOG, for paintings it may turn out that a size relating to the original size of the painting is a good way of getting around this problem.

C-HOG descriptors become more complex still. They are similar to Shape Context [2], but differ in one key aspect: in C-HOG descriptors each spatial cell holds a stack of gradient-weighted orientation cells over an orientation-independent edge-presence count which Shape Contexts use.

According to the author it is better to think of C-HOG descriptors as an advanced form of centre-surround coding as small descriptors with very few radial bins gave the best results.

Local contrast normalisation can be performed to help against local variations in the illumination of foreground and background.

It would seem that both R-HOG and C-HOG descriptors are designed more for the detection window rather than analysis technique. This may make them less useful and result in an implemented technique being just a simple histogram of edge orientations.

1.3.2 Brush-stroke Analysis

Stroke analysis is one of the main goals for this project. It is quite apparent from looking at Kyffin Williams' paintings that his brush-strokes change over time, his early work having lots of smaller strokes over the canvas to large bold strokes in his later work.

The first paper I found relating to the analysis of brush-strokes involved moving a circular filter across the whole painting to find the ridges of strokes, then filling any unbroken areas. They then shrunk these areas to a single pixel line and fitted a n^{th} order polynomial to this line [3]. This method seems fairly simplistic, but could be an interesting first step, but as it is more focused on authenticating paintings it may be of limited use.

Another method for stroke analysis has been published in the IEEE Transactions on Pattern Analysis and Machine Learning journal. This method is far more complex, but is able to extract and label individual brush-strokes. An interesting part of their findings was the ability to date some of Van Gogh's paintings to a known period in his career [15].

This method involves performing edge detection of the painting followed by an edge linking algorithm which aims to remove small, noisy edges and to trace every edge. With this they then perform enclosing, as strokes may not be complete this stage also aims to fill in missing gaps of strokes and to fill these in within a certain tolerance.

The algorithm then decides if a stroke really is a painted stroke, if the stroke is completely enclosed, isolated from other non-edge pixels and forms a connected component then it is likely that it is a proper brush-stroke and is extracted. The edge pixels are used as the background and the non-edge pixels as the foreground, this is the process of labelling the brush-stroke.

For each of these labelled candidates, a heuristic function is used to threshold any brush-strokes that are either too long or too short, these strokes are discarded. These strokes are then considered to be candidates if they are not significantly branched, the stroke is not too wide (this may change for Kyffin Williams as he used a pallet knife rather than a brush) and the brush-stroke is not too big or small.

Separately, the image is then segmented using k -means clustering by RGB values. This clustering algorithm is applied several times, lowering the tolerances for distance within a cluster. Connected components as a result of this clustering and have noise reduction performed upon them. Finally, the two types of brush-strokes are combined.

This technique may need some changing to account for Kyffin Williams' use of a pallet knife, but the overall principals of this technique should work with Kyffin's paintings.

1.4 Analysis Objectives

Analysis is one of the biggest sections of this project and involves creating techniques which will allow comparison of paintings in a way which will allow some form of classification to be performed on them.

Typically I would expect this to produce some form of high-dimension state space in which each painting is a point in the state space. From this state space the distance between one painting and another can be easily resolved using a distance measure like Manhattan distance (1.1), euclidean distance (1.2) or a distance measure more specific to the state space should it be needed (e.g.: chi-squared for histograms).

$$d_1(\mathbf{p}, \mathbf{q}) = \|\mathbf{p} - \mathbf{q}\|_1 = \sum_{i=0}^n |p_i - q_i| \quad (1.1)$$

$$d_1(\mathbf{p}, \mathbf{q}) = \sqrt{\sum_{i=0}^n (q_i - p_i)^2} \quad (1.2)$$

1.4.1 Colour-space Analysis

The simplest way of analysing a digital image is to look at the colours which it consists of. Doing this is relatively simple; each pixel has a set of values defining the colour of that point, getting something meaningful from this is less simple.

The simplest strategy is to perform some form of statistical analysis on each painting then use this for classification. Several good and computationally cheap options exist for this; mean (1.3) and standard deviation (1.4), are some good examples which often come predefined in image processing and computer vision libraries.

$$\mu = \frac{1}{N} \sum_{i=1}^N x_i \quad (1.3)$$

$$\sigma = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2} \quad (1.4)$$

The representation of colour is another important factor, an RGB representation will have all three values change if there are many changes in brightness of the colours whilst a Hue, Saturation, Value (HSV) representation will only have a single value change.

Therefore, an object of this section should be to explore different colour models and statistical methods which can be applied to them.

Another useful technique which should be investigated early into the project are image histograms. These histograms plot the distribution of colour across an image and are therefore a very powerful method of analysing an image, especially for comparison. As with statistical analysis, histograms will be largely effective by colour model.

1.4.2 Texture Analysis

As Kyffin Williams' work is very textural, it follows that a main part of the analysis should focus around the texture of his paintings. Unfortunately for this section, it seems unlikely that I will be able to get any 3-dimensional models of Kyffin's paintings. This would have been a nice, if rather large, section of the project.

Instead it is more sensible to look at the orientation of edges in Kyffin's work. Some useful pre-existing techniques have already been discussed in section 1.3.1. Histograms of edge orientation [9] seem like a promising concept which may prove relatively simple to implement.

This section may also help with any work into brush-stroke analysis (see section 1.4.3).

1.4.3 Brush-stroke Analysis

With Kyffin's distinctive style and how obviously this style changes over time, the ultimate aim of this project is to be able to analyse the brush-strokes¹ in a painting.

From looking at the paintings it is very apparent that in his earlier work he made a lot more strokes than in his later works². The strokes in his later work tend to have larger areas and span more of the canvas.

If it is possible to calculate a rough amount and size of strokes made in a given painting it should be a reasonable piece of data to classify on. As previously discussed in section 1.3.2 there has already been a decent amount of research into determining brush-strokes in a painting.

It would be preferable to try and take one of the techniques discussed in that research and change it to suit the needs of the project rather than attempting to create a whole new method of brush-stroke recognition.

1.4.4 Ensemble Techniques

With some of the aforementioned analysis techniques it makes sense to combine two or more techniques together; a good example would be colour histograms and histograms of edge orientation.

This form of analysis is inspired by the concept of the same name in statistics and machine learning which tend to obtain better predictive performance. It may also be worth while trying to weight different techniques so that the techniques which give the best performance affect the result of the ensemble technique more.

¹A slight misnomer as Kyffin used a palette knife to paint with rather than a traditional brush

²Although this isn't quite true as the canvases he worked on in his later life tended to be larger

1.5 Classification Objectives

The overall objective of classification is to be able to label a painting by Kyffin Williams as being painted in a given year based on analysis performed on all other paintings with known years.

This ties in with the main aim of this project of being able to classify any Kyffin Williams painting, whether it has a known or unknown year, as being from a given year. Evidently for paintings with an unknown year it is difficult to know how accurately the system has been, so, for the most part, these paintings have been ignored and those paintings with a known year have made up the training and validation set.

Because of the small size of paintings with known years it should be computationally viable to perform leave-one-out cross validation (figure 1.3).

```

function LOOCV(data)                                     ▷ data is a set of all data points
  for all item ∈ data do
    classifieditem ← CLASSIFY(item, data \ {item})
  end for
return classified
end function

```

Figure 1.3: Pseudocode for Leave-One-Out Cross Validation

This can be used to evaluate the performance of the analysis technique and classification algorithm. Pearson's product-moment correlation coefficient (1.5) between actual year and classified year has been suggested to be a good performance measure for this project.

$$\rho_{X,Y} = \frac{\text{cov}(X,Y)}{\sigma_X \sigma_Y} = \frac{E[(X - \mu_X)(Y - \mu_Y)]}{\sigma_X \sigma_Y} \quad (1.5)$$

1.5.1 Classification

One of the simplest methods of classification is k -Nearest Neighbour (figure 1.4) from this one can take a poll of the years for each neighbour and assign the year of the painting to classify to be the average of these years.

Depending which form of average you take (mathematical mean (1.3), median or mode) will alter the result; although it should be noted that median is very unlikely to give a result on its own due to the sparseness of the data.

```

Require:  $0 < k \leq |data|$                                      ▷ data is a set of all data points
function KNEARESTNEIGHBOUR(k, data)
  for  $i = 1 \rightarrow k$  do
    nni ← NEAREST(data)
    data ← data \ {nni}
     $i = i + 1$ 
  end for
return nn
end function

```

Figure 1.4: Pseudocode for k -Nearest Neighbour

There are other techniques which could be applied to this problem, but the rewards for implementing them is not likely to be outweighed by the time it would take to implement such

techniques. There is a workaround for this; there are several machine learning tool-kits which provide pre-implemented version of these techniques.

One of the most popular tool-kits available for general use is Weka [12], which is discussed in more detail in section 1.5.1.1.

Another technique suggested by Julie Greensmith is to use Learning Classifier Systems (LCS) [1], which has an implementation for Weka. This may prove to give very good results for the kind of analysis being performed on Kyffin's work.

1.5.1.1 Use of Weka

1.5.2 Exemplars

The use of exemplar images would be another way of performing classification. The idea of an exemplar is that a painting is the most representative of a given time period. With the help of Llyod, Lorna and the NLW a list of exemplars which can be used as a ground truth to classify against has been produced.

The initial idea for digitally producing exemplars is to take the middle painting for a time period, as would be expected. These can then be compared to the ground truths to see how correct the analysis technique performed.

However, there is also the potentially to generate a theoretical exemplar from the analysis. This might be hard to perform validation against the ground truth upon, but will give some useful data on Kyffin Williams' style and how it changed.

These theoretical exemplars would likely be produced using some form of Gaussian mixture model.

Chapter 2

Development Process

2.1 Introduction

2.2 Modifications

Chapter 3

Design

3.1 Methodology

Within the database of 325 paintings, the actual year of painting is known for 102 artworks. To judge how well an analysis technique performs I have used a leave-one-out cross validation methodology, only working with these 102 artworks so that any amount of uncertainty is removed.

Leave-one-out cross validation involves taking a painting for which the year is known, then use the classifier to guess the year; this allows the result to be validated against the known year, thereby showing how well a given analysis technique works.

To simplify the classification stage a k -Nearest Neighbour classifier is used with the remaining 101 paintings for which we also know the date of. k -NN is a fast, non-parametric classifier which is easy to implement and makes no assumptions about the underlying patterns in the data, merely the other paintings which are similarly located in the feature space which should be similarly located to paintings from around the same period of time is the analysis technique is performing well.

Figure 3.1 give a pictorial view of this process.

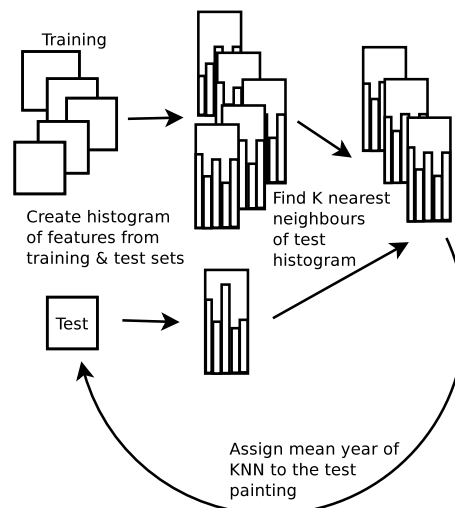


Figure 3.1: Overview of the Classification Methodology

3.2 Overall Architecture

The basic architecture for any system like this is to load the data in from a source of some form, apply an analysis technique to each data point then pass this data into the classification system.

From the classification system you should then be able to get the classified and actual year for each data point which can then have validation performed on it. This architecture is summed up in figure 3.2.



Figure 3.2: Basic Overall Architecture

Building up from this it is apparent that to implement the analysis and classification steps that there is a need to implement the factory method design pattern [11, p. 107-117]. Reading from a data source should be a simple matter of reading from a file, and cross validation has already been decided to use leave-one-out cross validation.

Figure 3.3 shows the design after adding in the factory methods.




Figure 3.3: Overall Architecture with Factory Methods

From this we then need the two top-level interfaces `Analyser` and `Classifier`. The `Analyser` interface should have a single method which runs analysis on a painting and return some form of object which represents the analysed data.

The `Classifier` class should have a method which takes a single painting and a set of paintings, returning a year which is the classified year of the single painting based on the set of paintings.

At this point it is also required that there is a class to store meta-data of a painting. Figure 3.4 depicts the design after adding in these parts.




Figure 3.4: Overall architecture with Interfaces for Analysers and Classifiers

Chapter 4

Implementation

4.1 Basic Structure

4.1.1 Loading Data

One of the more key parts to implement before all others in this project is the ability to load in data from Llyods spreadsheet. The first step of this was to convert it to a Comma-Separated Values (CSV) file-type, which is easier to read programatically.

The initial spreadsheet was slightly different from the version depicted below, Llyod was kind enough to update it so it was easier to handle digitally. In this version of the spreadsheet the filenames were much longer and were in sub-folders depending on their collection. Extra logic was needed to locate the file (a simple matter of concatenating the collection to the filename as a directory). The image files in the second version were just flat files which I decided would be better placed in a data directory.

The move to the newer version was a good excuse to clean up some of the initial code to use better Python programming practises; replacing messy loops with list comprehension where possible, using Key Word Arguments (`**kwargs`) and dictionaries instead of having utility methods, etc.

From the original format show in table 4.1 the data was converted to the CSV format show in figure 4.1.

Filename	ID	Title	Catalogue entry BBC YP	Genre	Height	Width	Area	Materials	Collection	image width	image height	image height/image width
001.jpg	1	A Chapel in the Tyrol	1950-1960	Landscape	40.7	29.8	1212.86	oil on hardboard	NLW	687	944	1.3741

Table 4.1: Layout of the Painting Data Spreadsheet

```
001.jpg,1,A Chapel in the Tyrol,1950-1960,Landscape,40.7,29.8,1212.86,oil on hardboard, NLW,687,944,1.3741
```

Figure 4.1: Painting Data in CSV Format

This CSV was then parsed using the Python 2.7 in-built `csv` module with relative ease. The example code (see listing D.1) helped with correct resource management as, at this point, my Python knowledge was fairly low.

4.1.2 Top-Level Classes

Being used to Java, an initial instinct of mine was to set up interfaces for the majority of top-level classes (namely the `Analyser` and `Classifier` classes). However, because of Python's

duck typing, there wasn't possibly to create these classes as interfaces. I decided to create them as abstract stub classes which just held placeholder method definitions which needed to be overridden in the concrete sub-classes. This was a slightly pointless exercise due to the weak typing, but useful for me to get my head around the architecture.

I also implemented the `Painting` class and had a method to take an array (from the CSV file) and create a new instance of it from this. Later this was changed to use Python `**kwargs` instead as it made more sense for the job. This change was done during the change to the second version of Llyods database.

4.1.3 Command Line Interface

As a research program with lots of different analysis and classification techniques the next step in setting up the basic architecture was to create a set of command-line arguments which would switch the functionality of the program. These arguments are depicted in table 4.2.

Name	Short Flag	Long Flag	Description
Analyser	-a	--analysers	Switch the analysis technique
Machine Learning	-m	--ml	Switch the machine learning technique
Data	-d	--csv	The data file to use
GUI	-g	--gui	Switch the GUI visualisation
Binning	-5	--bin-years	Put paintings in 5 year bins
Export	-e	--export	Export analysed data

Table 4.2: Command Line Arguments

The `argparse` library handles this nicely, the example code shown in listing D.2 shows the usage of this library.

These flags then hook into the factory methods to actually create the instances.

4.2 Colour Space Analysis

Colour space analysis involves performing statistical analysis on different colour models (RGB, HSV, etc.). This gives a very simplistic view of the entire image.

OpenCV offers the `Avg` method to perform the average across the image, however with a further look into the documentation there is also the `AvgStd` method which performs both mean and standard deviation on an image.

The analysed data was just the tuple returned by the `AvgStd` method. The distance measure was defined to be the sum of all elements in the tuple (in the case of an RGB colour model the mean red, green and blue and the standard deviation of red, green and blue).

4.2.1 Colour Models

There are many colour models to consider with digital image processing. RGB is one of the better known colour spaces as it is often how images are captured. It does have a problem in that all three values can change when the brightness changes.

As one of the main principals of this project is that Kyffin Williams' work darkened over time, it should follow that RGB may not be the best colour model to use.

To account for this it was decided to also use a HSV colour model to compare and contrast to RGB.

OpenCV handles colour spaces slightly oddly. Initially it uses `LoadImageM` to load the image, which uses an integer argument as a flag to define whether the image should be loaded in colour or grayscale.

From this image you then can use `CvtColor` to convert the colour model of an image, which uses an integer argument as a flag to define a number of different colour spaces.

Once converted, all methods act exactly the same as they would on a RGB image.

4.2.2 Colour Histograms

Colour Histograms are a representation of the distribution of colour across an image; this makes them very powerful for analysing the colour space of an image.

OpenCV provides a number of methods for both calculating and operating upon colour histograms. Here the example code (figure D.4) from the OpenCV documentation was a useful reference as some of the details of creating histograms in OpenCV are not noted implicitly in the python documentation.

The distance measure between colour histograms is also handled by OpenCV using a compare histogram method. Under the covers this uses a Chi Squared (equation 4.1) method to compare to histograms of equal dimensions.

$$d(H_1, H_2) = \sum_I \frac{(H_1(I) - H_2(I))^2}{H_1(I)} \quad (4.1)$$

4.3 Texture Analysis

4.3.1 Edge Count

The simplest technique to analyse texture is to apply an edge detection algorithm over the image and produce a count of all the marked edges regardless of orientation.

The Canny edge detector [6] provided a good output for measuring this, operators like Sobel provided less useful results and thus provided a worse analysis technique.

From the generated image a histogram was produced with only 2 bins; one bin for the presence of an edge, the other for the absence of an edge.

Open CV provides methods for applying the Canny edge detector to an image so this technique was simple to implement.

4.3.2 Edge Orientation

Edge orientation involves passing a filter over an image which is used to find the orientation of that gradient. There are many forms of filter which can be used to do this. The simplest of which discrete derivative masks (figure 1.1, figure 1.2, etc.), steerable filters are a more adjustable implementation of this. There are also more complex filters, like Gabor filters, which provide better flexibility and matching.

Using discrete derivative masks it is fairly simple to work out to orientation of a gradient mathematically. First pass a mask of the form $\begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}$, followed by one in the form $\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$.

These give you $\frac{\delta f}{\delta x}$ and $\frac{\delta f}{\delta y}$ respectively.

From these values you can then use a gradient direction function, shown in equation 4.2, to work out the actual direction of the gradient.

$$\theta = \text{atan2} \left(\frac{\delta f}{\delta x}, \frac{\delta f}{\delta y} \right) \quad (4.2)$$

Once this has been calculated a histogram of these orientations is created, this resulting histogram is then used as the analysed data. More complex operations can be performed on this, e.g. normalisation, which can potentially improve the overall result.

4.3.3 Steerable Filters

Another way of doing this is to change the orientation of the filter and bin the direction into certain cardinal directions; typically 0 , $\frac{\pi}{4}$, $\frac{\pi}{2}$ and $\frac{3\pi}{4}$ (from π to 2π becomes a repeat of any of those directions, only in reverse and are, therefore, covered already).

To do this we can use steerable filters $S(\theta)$ (shown in figure 4.2). With these we adjust the angle they are designed to match (θ), which will then give the degree to which a gradient matches that angle. To complete the binning by orientation, this value will then have a threshold applied, then added to the bin if the value is above the threshold.

$$S(0) = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 1 & 0 \\ 0 & 1 & 0 \end{bmatrix} S\left(\frac{\pi}{4}\right) = \begin{bmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{bmatrix} S\left(\frac{\pi}{2}\right) = \begin{bmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 0 & 0 & 0 \end{bmatrix} S\left(\frac{3\pi}{4}\right) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Figure 4.2: Steerable Filters

These have the advantage of being far less complex to compute than edge orientation; the maths and histogram creation on a floating point array are long running operations, to the point where canning the results for edge orientation is desirable.

There is the obvious trade off that steerable filters can only work in four directions which does affect the performance quite severely.

4.3.4 Gabor Filters

To increase the orientations matched from steerable filters without the need for increasingly large kernels, Gabor filters [?] were implemented. Gabor filters allow for an almost-infinite range of orientations without affecting the size of the filter. Equation 4.3 shows the mathematics behind Gabor Filters.

$$g(x, y; \lambda, \theta, \psi, \sigma, \gamma) = \exp \left(-\frac{x'^2 + \gamma^2 y'^2}{2\sigma^2} \right) \cos \left(2\pi \frac{x'}{\lambda} + \psi \right) \quad (4.3)$$

where:

$$\begin{aligned} x' &= x \cos \theta + y \sin \theta \\ y' &= x \sin \theta + y \cos \theta \end{aligned}$$

4.3.4.1 Histogram of Edge Orientation

4.4 Brush-Stroke Analysis

4.5 Classification and Validation

4.5.1 k -Nearest Neighbour

4.5.2 Leave-One-Out Cross Validation

4.5.3 Weka 3

4.5.3.1 Attribute-Relation File Format (ARFF)

4.5.4 Exemplars

4.5.4.1 Nearest Exemplar Classification

To implement Nearest Exemplar Classification was a fairly easy task: Llyod (with help from members of the NLW) provided a secondary spreadsheet which contained all the necessary information of exemplar by year (see table E.1 for the full document).

The spreadsheet was arranged in the format described in table 4.3, from there it was a simple matter of saving the spreadsheet as a CSV file and taking some of the existing code for parsing CSV files. This caused a slight problem in that the parsed data didn't have enough information to create a full `Painting` object, yet all the analysis techniques worked from these objects.

Filename	ID	Title	Catalogue Entry	Year
154.jpg	154	Landscape at Llanaelhaearn	1947	1947
<i>etc.</i>				

Table 4.3: Layout of the Exemplar Spreadsheet

This was solved easily thanks to Python's dynamic typing. A simple class which implemented all the necessary elements of `Painting` could be passed to the analysis techniques without any complaints. With a statically typed language this would have been harder to complete, but there would have been ways around using sub-classes and so on.

With the exemplars loaded and analysed, the program could continue as normal, until the classification step.

The idea of Nearest Exemplar Classification is to classify the unknown example using the nearest exemplar to that example in the feature space. This acts as a k -Nearest Neighbour with $k = 1$ and the space of neighbours only including the exemplars, rather than every other example. The pseudocode for this is shown in figure 4.3.

Initially this was implemented so that the examples that were exemplars were also classified, but this is a pointless exercise which only skews the results. Additional logic was added to skip any example which was an exemplar itself.

This proved to give slightly worse correlation per technique than k -Nearest Neighbour. This result is to be expected; for a start an artistically classified exemplar is unlikely to be the same as a statistically classified exemplar (see section 4.5.5). Also, with fewer examples to classify against, any variance in the data set (of which there is a lot) will likely be magnified.

```

function NEARESTEXEMPLARCLASSIFICATION(example, exemplars)
  if example  $\in$  exemplars then
    return NEAREST(example, exemplars \ example)
  else
    return NEAREST(example, exemplars)
  end if
end function

```

Figure 4.3: Nearest Exemplar Classification Psuedocode

Lastly, a painting may be picked as an exemplar by an expert for different reasons than any analysis technique that currently exists can give; emotional connections and knowledge of the artists history can be very subjective and may not relate to anything put down in paint.

4.5.5 Statistically Classified Exemplars

Another approach to exemplars is to work out a theoretical exemplar for a given period; the centroid of paintings within the given feature space for a single year, for example.

The simplest way of working this out is showing in figure 4.4, this works by taking the mean (equation 1.3) of each feature in the set of paintings for a single year, this will give the point in feature space that is most central. This is the same technique used to generate centroids in a clustering algorithm (*k*-Means Clustering, for example).

```

function STATISTICALCLASSIFYEXEMPLAR(examples)
  for all example  $\in$  examples do
    for all feature  $\in$  examplefeatures do
      averagefeature  $\leftarrow$  averagefeature + examplefeature
    end for
  end for
  for all feature  $\in$  average do
    averagefeature  $\leftarrow$   $\frac{\text{average}_{\text{feature}}}{\text{LENGTH}(\text{examples})}$ 
  end for
  return average
end function

```

Figure 4.4: Psuedocode for generating a Statistically Classified Exemplar

This may become a long operation depending on how many dimensions the feature space has. A technique like Principal Component Analysis (PCA) may be useful to help cut down the number of dimensions needed that this algorithm uses.

4.5.6 cv2 and Histogram Results

Whilst reading through the *numpy* [14] and *OpenCV* it turned out that, until this point, I had been using a deprecated version of *OpenCV*. *cv2* is the new version of *OpenCV* for python, which provides bindings for *OpenCV* 2.4 and uses *numpy* arrays and multi-dimensional arrays to represent all data types rather than in-built data structures.

This makes everything a lot easier to work with as all *numpy* arrays (with the correct data type) can be passed into the *cv2* methods.

4.6 3rd Party Libraries and Tools

4.6.1 Python

4.6.1.1 Python setuptools

4.6.2 OpenCV

4.6.3 scipy & numpy

4.6.4 matplotlib

4.6.5 Weka 3

4.6.5.1 liac-arff

4.6.6 git & github

Chapter 5

Testing

5.1 Overall Approach to Testing

As this is a research-based project, it is difficult to perform any for of unit, functional or requirements testing. Most testing is based on validating the results of analysis and classification techniques.

5.2 Validation

5.2.1 Leave-One-Out Cross Validation

5.2.2 Validation using Weka

Chapter 6

Evaluation

6.1 Evaluation of Requirements

6.2 Evaluation of Design

6.3 Evaluation of Tools

6.3.1 Programming Language

6.3.1.1 Dependency Management

6.3.2 Image Processing/Computer Vision Libraries

6.3.3 Machine Learning Libraries

6.3.3.1 File Formats

6.3.4 Scientific and Numeric Libraries

Appendices

Appendix A

Paper Submission for ISPA 2013

The following paper entitled *Can we date an artist's work from catalogue photos?* [5] has been submitted for the 8th International Symposium on Image and Signal Processing and Analysis.

Can we date an artist's work from catalogue photographs?

Alexander David Brown
Computer Science,
Aberystwyth University,
Penglais,
Aberystwyth,
Ceredigion,
Wales SY23 3DB
adb9@aber.ac.uk

Gareth Lloyd Roderick
School of Art,
Aberystwyth University,
Buarth Mawr,
Aberystwyth,
Wales SY23 1NG
glr7@aber.ac.uk

Hannah M. Dee
Computer Science,
Aberystwyth University,
Penglais,
Aberystwyth,
Ceredigion,
Wales SY23 3DB
hmd1@aber.ac.uk

Lorna M. Hughes
The National Library of Wales,
Aberystwyth,
Ceredigion,
Wales SY23 3BU
lorna.hughes@llgc.org.uk

Abstract—Kyffin Williams, art changes over time, blah blah blah. Features, colour, edges, histograms of oriented gradients; strong correlation using leave-one-out methodology. Exemplars; artistic and statistic.

I. INTRODUCTION

This paper presents a interdisciplinary computational study into the modelling of artistic style, and how this style changes over time. Sir John Kyffin Williams (1918-2006) was one of the predominant figures in Welsh art of the twentieth century. Kyffin – as he was almost universally known in Wales – studied at the Slade School of Art and worked as an art master at Highgate School, before returning to live on his native Anglesey in 1973. He was a prolific painter and once claimed to have painted “two pictures per week when in London, and three per week when in Wales.”[1, p.209] With a career spanning from the mid-1940s to approximately 2004, this rate amounts to a large body of work.

His style evolved from a very representational style to something more expressive, which retained representational qualities: the computer scientists on our team would say that the paintings became more *blocky*; the art historians that his landscapes are almost constructed with swathes of textural paint. His was a style characterised by thick impasto paint, applied almost exclusively with palette knife, although the application technique appears to change over time. This development of style led us to wonder: is it possible to date the pictures from images alone?

Through a collection of digital photographs of oil paintings, collected from museum websites, catalogues and other sources, we first investigate whether it is possible to date a painting based upon image features. We show that using a K-nearest neighbour classifier, tested using a leave-one-out methodology, we can obtain a strong correlation between image feature descriptors and year of painting. We go on to investigate whether exemplar based methods are able to improve on this, using what we call *artistic exemplars* (paintings selected by an expert as being typical for a particular year) and *statistic exemplars* (paintings which are near the centre of year-based clusters in feature space).

II. BACKGROUND

Although also a portrait painter, Williams is primarily known for his landscape paintings of north west Wales and Anglesey. While his technique and style changed over the years, his landscapes in oil are instantly recognisable, often featuring bold chunks of colour, and various points during his career bold black outlines to figures and landscapes features. Greens, browns and greys often form the palette of his paintings of the Welsh landscape. These colour selections seem appropriate for the artists claim that melancholy, derived from the “dark hills, heavy clouds and enveloping sea mists”, is a national characteristic of the Welsh.[1] This combination of colour selection and technique seems appropriate for the depiction of the areas where he painted. Many of his most successful paintings are said to have a “dark quality” in depicting “rain lashed hillsides,” and it was this darkness which “makes his landscapes so distinctively Welsh.”[2]

The aesthetic of Williamss Welsh landscapes is contrasted by the paintings he made following a trip to Patagonia to paint the landscape and people of the Welsh communities there in 1968 as part of a Winston Churchill Foundation scholarship. The colours and application of paint in pictures produced in following this journey (such as *Lle Cul*, *Henry Roberts*, *Bryngwyn Patagonia*, *Euros Hughes Irrigating his Fields*, all 1969, National Library of Wales) differ starkly from paintings of Welsh landscapes, incorporating pinks, purples and oranges. This contrast, combined with the fact that the Patagonian pictures were produced during a definite period of time has reinforced our interest in the analysis of the formal qualities of pictures from different collections remotely, using digital images.

Williamss work is well represented in public collections in Wales (particularly at the National Library of Wales, the National Museums and Galleries of Wales and Oriel Ynys Môn, Anglesey). His pictures, often depicting the landscape and people of north-west Wales were also tremendously popular with the art buying public. Of the 325 paintings by Williams in public collections in the UK listed on the BBC/Public Catalogue Foundation’s “Your Paintings” website, 212 of them are in the collections of the National Library of Wales[3]. Many of these paintings were bequeathed to the Library as

part of a larger bequest by the artist (including works on paper and other archival material). Many of the pictures which came to the library from the artists studio had little in the way of metadata, and as such have been catalogued with large date-ranges estimating the dates of production. This uncertainty in metadata is another motivating force behind the current project.

A. Taking a digital humanities approach to art history

Digital humanities is an established area of research that brings together digital content, tools and methods in order to address and create new knowledge across the disciplines. Digital humanities approaches can be seen in two distinct types of inquiry. The first is to carry out *traditional* humanities research more effectively or efficiently, by applying computational methods or approaches to digitized humanities sources (originally text, image, or audio-visual content from archives or libraries). Using John Unsworths definition of “scholarly primitives”[4] digital humanities scholarship customarily involves the use of digital tools and methods for discovering, annotating, comparing, referring, sampling, illustrating, or representing humanities data. A classic example of this sort of work would be the use of concordances and other computer-based analysis of digitized primary sources that have been processed by optical character recognition software to count, classify, or interpret digital texts (see, for example, the Historical Concordance of the Welsh language[5]). The second strand of digital humanities inquiry is the development of new research questions that can only be developed through the synthesis of digital content, tools and methods: work that would have otherwise been unimaginable[6]. This type of research is by necessity multi-disciplinary, drawing together expertise to be found across humanities, scientific and engineering disciplines, as well as involving content experts from libraries, archives and museums. However, in order to be truly transformative, this type of research must also be interdisciplinary.

The National Library of Wales now has a research programme in digital collections, which is a forum for investigation into the digital collections of Wales in collaboration with academics and students at universities in Wales and beyond, in order to develop new research based around the digital content created by the Library [7]. The research project described in this article is an example of a digital humanities collaborative venture, bringing together digital humanists, art historians, and computer scientists. The results of this research have value across all these groups. Arts historians are able to better investigate a large corpus of digital paintings through the application of computer science approaches to this content, and computer scientists are able to configure new approaches in imaging to working with a complex humanities data set.

B. Computer vision and the analysis of paintings

XXX Hannah to read [8], cite it sensibly, and pull some other references which are relevant

When we consider computer vision-based analysis of painterly style we find that the vast majority of work concentrates on brush stroke detection and analysis. For example, Berezhnoy and colleagues in [9] detect brush-strokes by moving a circular filter across the whole painting to find the ridges of strokes, then filling any unbroken areas. They then

shrunk these areas to a single pixel line and fitted a n^{th} order polynomial to this line.

Li et al [10] use a combination of edge analysis and clustering in colour space to determine strokes; a number of heuristics involving branching, stroke-width modeling, and gap filling are then used to refine the original brush stroke estimates. One interesting element of this work, from our perspective, is the ability to date some of Van Gogh’s paintings to a known period in his career.

Techniques based upon stroke analysis, whilst applicable to the work of some artists, are not applicable to all. In particular, Kyffin Williams painted with a palette knife and whilst there are clear *strokes* identifiable in his style, these vary widely in size and shape, so the morphological techniques which can detect strokes in Van Gogh’s work are unlikely to pay off when considering the blockier paintings in the Williams oeuvre. Another difference of note is that much work on computerised painting analysis (including [10], [9]) is based upon high resolution scans acquired in controlled conditions, whereas the current paper deals instead with a collection of photographs from catalogues, websites, and other disparate sources.

III. THE IMAGE DATASET

Our image dataset consists of 325 paintings, with associated metadata. Metadata includes title, year or year ranges (for those works where year is unknown but can be estimated by curators), genre, original painting size, painting materials and image size.

These photographs of paintings are challenging in and of themselves: they are not colour calibrated; some suffer from reflections (towards the end of his life Kyffin painted using exceptionally thick and textural strokes, which gives specularities on the catalogue images); they are at varying resolutions; and come from a range of different cameras. Image size bears little relation to the original painting size, and some images are even optimised for the web. Table I below summarises the dataset

Type	Number	Number (Known date)	Notes
Landscapes	247	64	
Portraits	52	35	
Seascapes	11	2	
Still lifes	4	1	
Other	8	0	Other or studies

TABLE I. A SUMMARY OF OUR KYFFIN WILLIAMS PAINTING DATASET

XXX It may be worth putting in something here about image size vs painting size?

IV. METHODOLOGY

Within our database of 325 paintings, we know the actual year of painting for 102 artworks. In order to determine the accuracy of our results, rather than work with the full dataset (and work with images with uncertain metadata in the form of date ranges), we have used a leave-one-out cross validation methodology. This involves us taking a painting for which we know the year, and then using our classifier to guess that year;




Fig. 1. Overview of the classification methodology

thus we are able to tell whether we are right. We are also able, if we are wrong, to determine exactly how wrong we are.

To simplify the classification stage we use a K-Nearest Neighbour (KNN) classifier with the other 101 paintings for which we know the date. KNN is a fast, non-parametric classifier which makes no assumptions about the underlying patterns in the data, merely that paintings from around the same time will be similarly located in our feature space(s). Whilst we suspect that there may be some broader underlying trend in the change of style, for this work have concentrated on features for classification rather than the question of classification or regression itself.

Thus for each feature set, we take all paintings for which we know the year of creation; select one painting, and find its nearest neighbours within that feature space. The year assigned by our classifier to that painting is the mean of the K neighbours; we found this provided better results than both median and mode. Figure 1 provides an overview of this classification methodology.

We also know that painting’s actual year, and we can plot actual against predicted year for all known-year paintings. To measure goodness of fit, the Pearson’s product-moment correlation coefficient was calculated on these orderings; this provides us with a performance measure of each classifier. It is also possible to test Pearson’s R for statistical significance; thus significance levels are reported alongside R in this paper. With all of the feature spaces we consider, it is possible treat the painting descriptors as histograms. This allows us to use a single distance measure, namely chi-squared, in our K-nearest neighbour classification.

V. AN EXPLORATION OF COLOUR AND TEXTURE FEATURES

The digital analysis of paintings is a broad research area. Within the methodology we have selected, there are many feature spaces which could be useful: from simple analysis of the way in which colour changes over time, through edge detection, to texture analysis. We have concentrated on lower level image features – colours, textures, and edges – rather than attempt to extract brush strokes. As mentioned earlier, Williams painted with a palette knife rather than a brush, and his work is characterised by angularity rather than identifiable “strokes”. In this section we describe the various feature sets and feature spaces we have explored; results for each of these are presented in Section VI below.

There is a clear (to the eye) trend in colour usage, as the paintings get “gloomier” over time, so we started with simple colour-space analysis: taking the mean RGB for each painting and using this with our KNN classifier; we also tested other colour spaces, such as HSV. Promisingly this provided us with a positive correlation.

Staying with the colour variation theme, we then used colour histograms, which provide a more precise representation of the way Kyffin Williams used colour. These histograms were developed by counting the number of pixels within a particular colour range for each painting, and then building a normalised histogram representing the colour usage.

As a lot of Kyffin Williams’ paintings are highly textural, edge detection and texture analysis was thought to be a good avenue to explore. Firstly, we investigated simple *edginess*; as a rough estimate of the the edge properties of the artworks we apply a Canny [11] edge detector to the paintings, and then use a count of edge pixels as our feature.

Texture analysis is a continuation of edge detection. Instead of taking simply the strength and number of edges, we create a histogram of orientated gradients as in [12]. In this way we begin to build up a richer representation of the texture of a painting. Given the change in style of Kyffin Williams’ work, moving away from figurative representations with curved lines towards more blocky rectilinear brush strokes, we expect these edge orientation frequencies to change over time. To this end we used simple steerable filters S , applied to the image at 0 , $\frac{\pi}{4}$, $\frac{\pi}{2}$ and $\frac{3\pi}{4}$.

$$S\left(\frac{\pi}{2}\right) = \begin{pmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 0 & 0 & 0 \end{pmatrix} \quad (1)$$

Equation 1 shows a sample steerable filter, in this case $S(\frac{\pi}{2})$, the filter which gives the highest response when presented with horizontal lines. By convolving each image with filters tuned to different orientations, we can build a histogram recording the frequency of lines at each orientation.

Gabor filters are linear filters which can be tuned to a greater range of angles and frequencies than simple steerable filters, which in turn results in a more accurate representation of the texture of the painting. The general equation for a Gabor filter is given in Equation 2.




Fig. 2. Example image – Coastal Sunset, 1990-2006




Fig. 3. Steerable filter strength $S(\theta)$ on the example image in figure 2

$$g_e(x) = \frac{1}{2\pi\sigma_x\sigma_y} e^{-\frac{1}{2}\left(\frac{x^2}{\sigma_x^2} + \frac{y^2}{\sigma_y^2}\right)} \cos(2\pi\omega_{x_0}x + 2\pi\omega_{y_0}y) \quad (2)$$

Where $(\omega_{x_0}, \omega_{y_0})$ defines the centre frequency, and (σ_x, σ_y) the spread of the Gaussian window. In this work we use Gabor filters tuned to equally spaced orientations to build a histogram representing line orientations in each painting, and present results below for histograms built from the output of both 4 and 8 filter orientations.

XXX Hannah to look at a textbook for a better version of the above scary maths and cite.

The final method for producing histograms we consider involves the application of two discrete derivative masks to the image to get the gradient of x and y , and then to work out the gradient direction at each point. These gradient directions are then summarised in a histogram of oriented gradients, providing a yet richer representation of the texture of the image. This is similar to the method described in [12].

XXX Alex we use 15 bins; can we try different bin size? or is that just me being fussy? It'd be good to have Gabor (4, 8, 16); HOG (4, 8, 16); thus providing a direct comparator

VI. YEAR CLASSIFICATION RESULTS

The one parameter of our classifier is the choice of K in K -nearest neighbour. Simply setting $K = 1$ has the effect of assigning the year of the nearest painting in feature space to the current test painting, whereas setting $K = 102$ has the




Fig. 4. Correlation Coefficients r against K values for K -Nearest Neighbour

effect of giving each painting the mean value of the entire dataset. Clearly a point between these two extremes would be best; from Figure VI we can see that for many of the feature spaces we consider, the optimum K value is around 7 or 8.

Technique	r	$P(r)$
Edge Strength	0.01072814442	0.910207365322
HSV	0.11224162320	0.236562323053
RGB	0.11775991536	0.214157639943
HSV Histograms	0.14603200925	0.122741593497
RGB Histograms	0.26953009140	0.00389039298815
Steerable filters: 4 orientations	0.31169031994	0.00077735274356
Gabor filters: 8 orientations	0.43790920108	1.23123557269e-06
Gabor filters: 4 orientations	0.44084541429	1.02469784947e-06
HOG (Discrete Derivatives): 15 orientations	0.44695489816	6.95435503135e-07

TABLE II. CORRELATION COEFFICIENTS, ORDERED BY STRENGTH

VII. EXEMPLARS: CAN WE IMPROVE RESULTS BY INCORPORATING EXPERT KNOWLEDGE?

We have also investigated the utility of incorporating expert knowledge within our framework. For each year represented in our collection we asked Dr Paul Joyner, of the National Library of Wales, to choose the one painting which best represents the artist's work for that year. Dr Joyner is a member of the Trustees of the Kyffin Williams Estate and he has written widely on Welsh Art and Kyffin Williams. These chosen paintings we consider to be connoisseurially/artistically selected exemplars (*artistic exemplars*, for short), which we can then use as a representation of that particular year.

The data on artistic exemplars opens up the options for different methods of classification. Rather than using K -nearest neighbour to classify each point in the feature space, we can

take the year of the nearest exemplar to assign that year to the painting in question.

This then raises the question of whether we can determine *statistical exemplars* to compare with our *artistic exemplars*, and if so what the digitally-chosen exemplars would be. We can either use cluster centres (which provides us with a point in feature space which will not correspond to an actual painting), or the nearest actual painting to the feature space centroid for a year. The former technique does not, strictly speaking, give us an exemplar; the latter chooses as exemplar the painting which best represents a particular year according to a particular feature space.

These exemplars give an interesting insight into the feature space, for example, the artistic exemplar *Landscape with Cattle* for 1950 is often correctly classified as the painting nearest to the statistical centroid.

In contrast, the artistic exemplar for 1960 - *Nant Ffrancon from Llandegfan* is commonly misclassified as *Moelwyn Bach* and is never correctly classified as the statistical exemplar.

XXX Alex - We need results:-)

XXX Alex again - can we get an indication of any situations where our statistic exemplars match the artistic ones, and the converse situation (artistic exemplars nowhere near centres of feature space)? This is a good place for us to include a couple of extra photos of paintings

A. Comparing Artistic and Statistical Exemplars

The same distance measure used to run k-nearest neighbour on the feature-based classifiers was used to generate the error between the statistical and artistic exemplars.

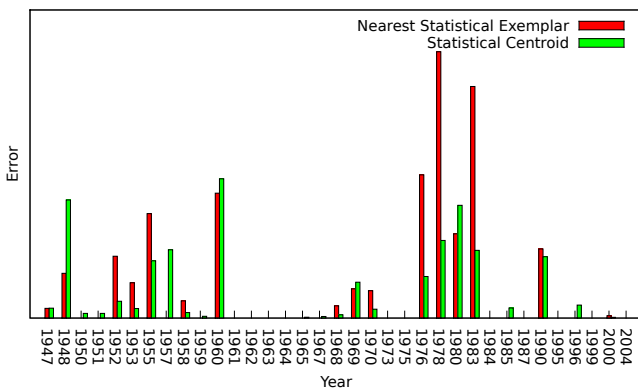


Fig. 5. Statistical Exemplar Error by Year for Gabor Filters (8 orientations)

XXX Alex... er... your to do list is longer than mine:-)

VIII. CONCLUSIONS AND FUTURE DIRECTIONS

To the best of our knowledge this is the first work that attempts to date work by an artist by year. Similarly, we believe we are the first to try and perform digital analysis of paintings from a range of catalogue and web images. The results presented here show that computer vision *can* help with the job of dating art within an artist's body of work.

Future directions will involve testing the methods presented here on the works of other artists who have shown great stylistic variation over the course of their career: we would like to build a dataset of, for example, David Hockney works. Whilst we have not yet performed this test we are hopeful of success; by avoiding brushstroke detection (which we expect to be artist specific) we hope to have developed techniques with application across a broader range of artistic styles.

ACKNOWLEDGMENT

The authors would like to thank Dr Paul Joyner of the National Library of Wales for his invaluable expert assistance. We would also like to thank Professor Robert Meyrick of the School of Art, Aberystwyth University.

REFERENCES

- [1] K. Williams, *Across the Straits: An Autobiography*. Llandysul, Wales: Gomer, 1993.
- [2] J. Davies, *100 Welsh heroes*. Aberystwyth, Wales: Culturenet Cymru, 2004.
- [3] "Your Paintings - Kyffin Williams," Mar. 2013. [Online]. Available: http://www.bbc.co.uk/arts/yourpaintings/paintings/search/painted_by/kyffin-williams_artists
- [4] J. Unsworth, "Scholarly primitives: what methods do humanities researchers have in common, and how might our tools reflect this?" King's College London, 2000. [Online]. Available: <http://www3.isrl.illinois.edu/~unsworth/Kings.5-00/primitives.html>
- [5] "Corpws Hanesyddol yr Iaith Gymraeg 1500-1850/A Historical Corpus of the Welsh Language 1500-1850," 2004, <http://people.ds.cam.ac.uk/dwew2/hcwl/menu.htm>.
- [6] L. M. Hughes, *Evaluating and Measuring the Value, Use and Impact of Digital Collections*. London: Facet, 2011.
- [7] <http://www.llgc.org.uk/research>.
- [8] D. G. Stork, "Computer vision and computer graphics analysis of paintings and drawings: An introduction to the literature," ser. Lecture Notes in Computer Science, X. Jiang and N. Petkov, Eds. Berlin, Heidelberg: Springer Berlin / Heidelberg, 2009, vol. 5702, ch. 2, pp. 9-24.
- [9] I. E. Bereznoy, E. O. Postma, and Herik, "Automatic extraction of brushstroke orientation from paintings," *Machine Vision and Applications*, vol. 20, no. 1, pp. 1-9, 2009.
- [10] J. Li, L. Yao, E. Hendriks, and J. Z. Wang, "Rhythmic brushstrokes distinguish van gogh from his contemporaries: Findings via automated brushstroke extraction," *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, vol. 34, no. 6, pp. 1159-1176, Jun. 2012.
- [11] J. F. Canny, "A computational approach to edge detection," *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, vol. PAMI-8, no. 6, pp. 679-698, Nov. 1986.
- [12] N. Dalal and B. Triggs, "Histograms of oriented gradients for human detection," in *Computer Vision and Pattern Recognition, 2005. CVPR 2005. IEEE Computer Society Conference on*, ser. CVPR '05, vol. 1. Washington, DC, USA: IEEE, Jun. 2005, pp. 886-893 vol. 1.

Appendix B

3rd Party Libraries and Tools

2.1 Python 2.7

2.1.1 setuptools

<http://pypi.python.org/pypi/setuptools>

2.1.2 scipy

<http://www.scipy.org/> [14]

2.1.3 numpy

<http://www.numpy.org/> [14]

2.1.4 matplotlib

<http://matplotlib.org/>

2.1.5 liac-arff

<https://github.com/renatopp/liac-arff>

2.2 OpenCV

<http://opencv.org/> [4]

2.2.1 OpenCV Python

<http://opencv.willowgarage.com/wiki/PythonInterface> [4]

2.3 Weka 3

<http://www.cs.waikato.ac.nz/ml/weka/> [12]

2.4 git

2.4.1 github

Appendix C

Equations

3.1 Statistical Equations

3.1.1 Mean

$$\mu = \frac{1}{N} \sum_{i=1}^N x_i$$

3.1.2 Standard Deviation

$$\sigma = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2}$$

3.1.3 Pearson's product-moment coefficient

$$\rho_{X,Y} = \frac{\text{cov}(X, Y)}{\sigma_X \sigma_Y} = \frac{E[(X - \mu_X)(Y - \mu_Y)]}{\sigma_X \sigma_Y}$$

3.2 Distance Equations

3.2.1 Manhattan Distance

$$d_1(\mathbf{p}, \mathbf{q}) = \|\mathbf{p} - \mathbf{q}\|_1 = \sum_{i=0}^n |p_i - q_i|$$

3.2.2 Euclidean Distance

$$d_1(\mathbf{p}, \mathbf{q}) = \sqrt{\sum_{i=0}^n (q_i - p_i)^2}$$

3.3 Filter Equations

3.3.1 Gradient Direction

$$\theta = \text{atan2} \left(\frac{\delta f}{\delta x}, \frac{\delta f}{\delta y} \right)$$

3.3.2 Discrete Derivative Masks

$$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}$$

3.3.3 Gabor Filter

$$g(x, y; \lambda, \theta, \psi, \sigma, \gamma) = \exp\left(-\frac{x'^2 + \gamma^2 y'^2}{2\sigma^2}\right) \cos\left(2\pi \frac{x'}{\lambda} + \psi\right) \quad (\text{C.1})$$

where:

$$\begin{aligned} x' &= x \cos \theta + y \sin \theta \\ y' &= x \sin \theta + y \cos \theta \end{aligned}$$

Appendix D

Code Samples

4.1 Example CSV Parsing Code

Listing D.1: Example CSV Parsing Code from <http://docs.python.org/2/library/csv.html>

```
>>> import csv
>>> with open('eggs.csv', 'rb') as csvfile:
...     spamreader = csv.reader(csvfile, delimiter=',',
...     quotechar='|')
...     for row in spamreader:
...         print ', '.join(row)
Spam, Spam, Spam, Spam, Spam, Baked Beans
Spam, Lovely Spam, Wonderful Spam
```

4.2 argparse Example Code

Listing D.2: Example argparse Code from <http://docs.python.org/2/library/argparse.html>

```
import argparse

parser = argparse.ArgumentParser(description='Process some integers.')
parser.add_argument('integers',
                    metavar='N',
                    type=int,
                    nargs='+',
                    help='an integer for the accumulator')
parser.add_argument('--sum',
                    dest='accumulate',
                    action='store_const',
                    const=sum, default=max,
                    help='sum the integers (default: find the max)')

args = parser.parse_args()
print args.accumulate(args.integers)
```

4.3 Gabor Filter Example Implementation

Listing D.3: Example implementation of a Gabor Filter in MATLAB from wikipedia [8]

```
function gb=gabor_fn(sigma,theta,lambda,psi,gamma)

sigma_x = sigma;
sigma_y = sigma/gamma;

% Bounding box
nstds = 3;
xmax = max(abs(nstds*sigma_x*cos(theta)),abs(nstds*sigma_y*sin(theta)));
xmax = ceil(max(1,xmax));
ymax = max(abs(nstds*sigma_x*sin(theta)),abs(nstds*sigma_y*cos(theta)));
ymax = ceil(max(1,ymax));
xmin = -xmax; ymin = -ymax;
[x,y] = meshgrid(xmin:xmax,ymin:ymax);

% Rotation
x_theta=x*cos(theta)+y*sin(theta);
y_theta=-x*sin(theta)+y*cos(theta);

gb= exp(-.5*(x_theta.^2/sigma_x^2+y_theta.^2/sigma_y^2)).*cos(2*
    pi/lambda*x_theta+psi);
```

4.4 OpenCV Histogram Example Code

Listing D.4: Example Histogram calculation and displaying code from OpenCV [4].

```
# Taken from: http://opencv.willowgarage.com/documentation/
python/imgproc-histograms.html#calchist
# Calculating and displaying 2D Hue-Saturation histogram of a
color image
```

```
import sys
import cv

def hs_histogram(src):
    # Convert to HSV
    hsv = cv.CreateImage(cv.GetSize(src), 8, 3)
    cv.CvtColor(src, hsv, cv.CV_BGR2HSV)

    # Extract the H and S planes
    h_plane = cv.CreateMat(src.rows, src.cols, cv.CV_8UC1)
    s_plane = cv.CreateMat(src.rows, src.cols, cv.CV_8UC1)
    cv.Split(hsv, h_plane, s_plane, None, None)
    planes = [h_plane, s_plane]

    h_bins = 30
```

```

s_bins = 32
hist_size = [h_bins, s_bins]
# hue varies from 0 (~0 deg red) to 180 (~360 deg red again
  */
h_ranges = [0, 180]
# saturation varies from 0 (black-gray-white) to
# 255 (pure spectrum color)
s_ranges = [0, 255]
ranges = [h_ranges, s_ranges]
scale = 10
hist = cv.CreateHist([h_bins, s_bins], cv.CV_HIST_ARRAY,
    ranges, 1)
cv.CalcHist([cv.GetImage(i) for i in planes], hist)
(_, max_value, _, _) = cv.GetMinMaxHistValue(hist)

hist_img = cv.CreateImage((h_bins*scale, s_bins*scale), 8,
    3)

for h in range(h_bins):
    for s in range(s_bins):
        bin_val = cv.QueryHistValue_2D(hist, h, s)
        intensity = cv.Round(bin_val * 255 / max_value)
        cv.Rectangle(hist_img,
            (h*scale, s*scale),
            ((h+1)*scale - 1, (s+1)*scale - 1),
            cv.RGB(intensity, intensity, intensity),
            cv.CV_FILLED)

    return hist_img

if __name__ == '__main__':
    src = cv.LoadImageM(sys.argv[1])
    cv.NamedWindow("Source", 1)
    cv.ShowImage("Source", src)

    cv.NamedWindow("H-S_Histogram", 1)
    cv.ShowImage("H-S_Histogram", hs_histogram(src))

    cv.WaitKey(0)]

```

Appendix E

Spreadsheet Data

Filename	ID	Title	Catalogue entry BBC YP	exemplar year
154.jpg	154	Landscape at Llanaelhaearn	1947	1947
258.jpg	258	Snowdon, the Traeth and the Frightened Horse	1948	1948
155.jpg	155	Landscape with Cattle	c.1950	1950
286.jpg	286	The Dark Lake	1951	1951
288.jpg	288	The Moelwyns from Aberglasyn	c.1952	1952
238.jpg	238	Self Portrait	c.1953	1953
051.jpg	51	Cottages, Llanddona	c.1955	1955
185.jpg	185	Mountain Landscape	1957	1957
023.jpg	23	Capel Carmel	c.1958	1958
246.jpg	246	Snow above Beddgelert	c.1959	1959
206.jpg	206	Nant Ffrancon from Llandegfan	c.1960	1960
260.jpg	260	Snowstorm off Penmon	1961	1961
053.jpg	53	Cottages, Mynydd Bodafon	c.1962	1962
120.jpg	120	German Girl	c.1963	1963
243.jpg	243	Sir David Hughes Parry	1964	1964
242.jpg	242	Sir Charles Evans	c.1965	1965
087.jpg	87	Farm below Crib Goch	1967	1967
248.jpg	248	Snow on Siabod	c.1968	1968
158.jpg	158	Lle Cul, Patagonia	1969	1969
073.jpg	73	Deserted Farm, Llanrhuddlad	c.1970	1970
275.jpg	275	Sun and Cloud on Lliwedd	1973	1973
093.jpg	93	Farm, Llanfairynghornwy	c.1975	1975
018.jpg	18	Blaen Nant	1976	1976
237.jpg	237	Sea at Trearddur	c.1976	1976
017.jpg	17	Blaen Ffrancon No.1	c.1978	1978
107.jpg	107	Farmers on the Carneddau	c.1980	1980
105.jpg	105	Farmer below the Ridge	c.1983	1983
202.jpg	202	Mrs Rowlands	c.1984	1984
003.jpg	3	Above Carneddi, No.2	c.1985	1985
264.jpg	264	Storm at Trearddur	1987	1987
165.jpg	165	Llyn-y-Cau, Cader Idris	c.1990	1990
267.jpg	267	Storm, Porth Cwyfan	1995	1995
268.jpg	268	Storm, Trearddur	1996	1996
214.jpg	214	Pengwryd	c.1999	1999
020.jpg	20	Bryn Cader Faner	2000	2000
278.jpg	278	Sunset, Anglesey	2004	2004

Table E.1: Exemplar Spreadsheet Data

Annotated Bibliography

- [1] J. Bacardit and X. Llorà, “Large-scale data mining using genetics-based machine learning,” *Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery*, vol. 3, no. 1, pp. 37–61, Jan. 2013. [Online]. Available: <http://dx.doi.org/10.1002/widm.1078>

A paper recommended by Julie Greensmith for information on a Learning Classifier System (LCS) which Julie believes will yield good results with the Kyffin Williams project.

- [2] S. J. Belongie, J. Malik, and J. Puzicha, “Matching shapes,” in *Computer Vision, 2001. ICCV 2001. Proceedings. Eighth IEEE International Conference on*, vol. 1. IEEE, 2001, pp. 454–461 vol.1. [Online]. Available: <http://dx.doi.org/10.1109/iccv.2001.937552>
- [3] I. E. Berezhnoy, E. O. Postma, and H. J. van den Herik, “Authentic: Computerized Brushstroke Analysis,” in *Multimedia and Expo, 2005. ICME 2005. IEEE International Conference on*. IEEE, July 2005, pp. 1586–1588. [Online]. Available: <http://dx.doi.org/10.1109/icme.2005.1521739>

Defines a method of analysing brushstrokes by applying a circular filter across a digital image to pick up the ridges of a brushstroke. This can then be used to pick out individual brushstrokes in order to be able to fit a nth order polynomial to them. Though this paper focuses on authenticating Van Gogh’s paintings, it could easily be applied to the work of Kyffin Williams and may allow for some interesting analysis.

- [4] G. Bradski, “The OpenCV Library,” *Dr. Dobb’s Journal of Software Tools*, 2000.

Used Python (<http://opencv.willowgarage.com/documentation/python>) and C++ (<http://opencv.willowgarage.com/documentation>) documentation for library reference and some learning on image processing/computer vision. Used since 11 October 2012.

- [5] A. D. Brown, H. M. Dee, G. L. Roderick, and L. M. Hughes, “Can we date an artist’s work from catalogue photographs?” Apr. 2013.

The paper submitted to ISPA as a part of this project.

- [6] J. F. Canny, “A Computational Approach to Edge Detection,” *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, vol. PAMI-8, no. 6, pp. 679–698, Nov. 1986. [Online]. Available: <http://dx.doi.org/10.1109/tpami.1986.4767851>

Reference for the Canny Edge detection technique.

- [7] I. Chilvers, J. Glaves-Smith, and I. Chilvers, *A dictionary of modern and contemporary art*. Oxford University Press, 2009. [Online]. Available: <http://www.worldcat.org/isbn/0199239665> 0199239665.

References for Sir John Kyffin Williams

- [8] W. Contributors, “Gabor filter,” Online, Oct. 2012. [Online]. Available: http://en.wikipedia.org/w/index.php?title=Gabor_filter&oldid=517342109

Reference for the example code for Gabor Filters. A better reference should be found if possible.

- [9] N. Dalal and B. Triggs, “Histograms of oriented gradients for human detection,” in *Computer Vision and Pattern Recognition, 2005. CVPR 2005. IEEE Computer Society Conference on*, ser. CVPR ’05, vol. 1. Washington, DC, USA: IEEE, June 2005, pp. 886–893 vol. 1. [Online]. Available: <http://dx.doi.org/10.1109/cvpr.2005.177>

Describes a method of producing histograms of edge orientations which may prove to be a useful analysis technique for Kyffin Williams’ art. The most interesting part of this paper is the use of segmentation and binning of gradients which seems like it could be useful to differentiate different parts of the image which may be painted in different styles.

- [10] J. Davies and A. Gymreig, *The Welsh Academy encyclopaedia of Wales*. University of Wales Press, 2008, pp. 957–958. [Online]. Available: <http://www.worldcat.org/isbn/9780708319536> 9780708319536.

References for Sir John Kyffin Williams.

- [11] E. Gamma, R. Helm, R. Johnson, and J. Vlissides, *Design Patterns: Elements of Reusable Object-Oriented Software*, ser. Addison-Wesley professional computing series. Addison-Wesley, 1996. [Online]. Available: <http://www.worldcat.org/isbn/9780201633610> 9780201633610.

Design Pattern reference for the architecture of the program.

- [12] M. Hall, E. Frank, G. Holmes, B. Pfahringer, P. Reutemann, and I. H. Witten, “The WEKA data mining software: an update,” *SIGKDD Explor. Newsl.*, vol. 11, no. 1, pp. 10–18, Nov. 2009. [Online]. Available: <http://dx.doi.org/10.1145/1656274.1656278>

Citation for the Weka data mining software. Weka is a Java based tool which can be used to run a lot of classifiers to a dataset, making it a very useful tool to apply to the Kyffin Williams project. Weka allows the application of complex machine learning techniques without having to spend a lot of time learning, understand and implementing said techniques.

- [13] R. Harris, “How Rolf learnt to paint like Sir Kyffin Williams,” BBC Broadcast, Feb. 2011. [Online]. Available: <http://www.bbc.co.uk/programmes/p00f6nyt>

A video on the BBC by Rolf Harris about some of Kyffin Williams’ life and about his interesting style of painting.

- [14] E. Jones, T. Oliphant, P. Peterson, *et al.*, “SciPy: Open source scientific tools for Python,” 2001.

Reference for SciPy and NumPy; two useful libraries which implement complex scientific and mathematical functions. NumPy is also heavily tied to OpenCV version 2.4 so it is a necessity to include as part of this project.

- [15] J. Li, L. Yao, E. Hendriks, and J. Z. Wang, “Rhythmic Brushstrokes Distinguish van Gogh from His Contemporaries: Findings via Automated Brushstroke Extraction,” *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, vol. 34, no. 6, pp. 1159–1176, June 2012. [Online]. Available: <http://dx.doi.org/10.1109/tpami.2011.203>

Defines a complex method for analysing individual brush strokes which has been used to classify the period of two paintings by Van Gogh. This technique could be very powerful when applied to Kyffin Williams’ work. This could be one of the most important techniques for the whole of the Kyffin Williams project.