

Kyffin Williams: Digital Analysis of Paintings

Report Name	Progress Report
Author (User Id)	Alexander D. Brown (adb9)
Supervisor (User Id)	Hannah M. Dee (hmd1)
Module	CS39440
Date	November 14, 2012
Revision	0.3
Status	Draft
Word Count	?

Contents

1	Project Summary	2
2	Current Progress	2
2.1	Library and Language Research	2
2.2	Other Tools	3
2.3	Design	6
2.4	Colour Space Statistical Analysis	6
2.5	Histogram Analysis	7
2.6	Nearest Neighbour Classification	7
2.7	Distance Measures	7
2.8	Meetings with Gareth Lloyd Roderick	7
2.9	Leave-one-out Cross Validation	7
2.10	Research into Stroke Analysis	9
2.11	Technical Issues	9
3	Planning	10
3.1	Methodology	10
3.2	Gantt Chart	12
3.3	Mid-term Demonstration	12
	Annotated Bibliography	13

List of Figures

1	Using OpenCV to blur an image	4
2	Using FIJI to blur an image	4
3	Using IVT to blur an image	5
4	Using VXT to blur an image	5
5	Initial Design	6
6	Mean RGB intensities by year.	8
7	Mean HSV intensities by year.	8
8	1-Nearest Neighbour	9
9	K-Nearest Neighbour	9
10	Manhattan Distance	10
11	Euclidean Distance	10
12	Leave-one-out cross validation of actual year against classified year using RGB statistical analysis and 1-nearest neighbour classification.	11

List of Tables

1	Comparison of image processing/computer vision libraries.	2
---	---	---

1 Project Summary

Sir John “Kyffin” Williams was a landscape painter from Wales whose work was predominantly based in Wales and Patagonia. Gareth Lloyd Roderick, with the help of the National Library of Wales, has collected data, such as the date or location, of these paintings. This data allows for some interesting analysis; particularly that of temporal or geological classification of a given painting. That is being able to take a painting and decide the year or location in which it was painted from a database of existing, known, works by Kyffin Williams.

Temporal analysis will be the focus of this project as it allows for a diverse range of techniques; from statistical analysis of colour values of the paintings to looking at the length and style of paintbrush strokes. Geological analysis would likely be very difficult, especially as the locations depicted were often sketched on-site then painted in a studio.

Whilst it would be nice to be able to predict the age of a painting with no known year, it is far more interesting to try to guess the year of paintings for which the date is known. This project will use leave-one-out cross-validation to help measure the effectiveness and validity of the analysis techniques employed in this project. Leave-one-out validation can be used with this project as the data set is small enough not to incur large performance overheads and the overall speed of the program is irrelevant so long as it completes within a decent amount of time.

One major limitation with this is it also includes the technique used for classification, some techniques might work better with K-Nearest Neighbour whilst other might benefit from more complex methods of classification.

2 Current Progress

2.1 Library and Language Research

The first step in the project was to research into the numerous image processing/computer vision libraries available for use with the aim to find the most suitable library to use constantly for the rest of the project. Because of this the library had to be easy to install and use, as well as having all the necessary features which would be used in this project.

The choice of library also has a knock-on affect on the choice of language this project would be written in. To this end I downloaded each of the libraries in turn, created a simple application to perform blur on a single image to test their use and documentation. From this I gained a good insight into how easy that library would be to perform more complex tasks and how confident I felt with the language(s) the library has bindings for.

Table 1 shows an overview of all the libraries I have currently considered and experimented with.

Library	Platform				Language(s)	Example
	Windows	Mac	Linux	Android		
OpenCV	✓	✓	✓	✓	C, C++, Python	Figure 1
FIJI	✓	✓	✓		Java	Figure 2
IVT	✓	✓	✓		C++	Figure 3
VXL	✓	✓	✓		C++	Figure 4

Table 1: Comparison of image processing/computer vision libraries.

Of these libraries OpenCV was the easiest to work with. OpenCV also boasts a wide range of features, all of which are well documented. FIJI provided a lot of high-level functionality, but for use as a library it quickly became unwieldy and was difficult to find the correct classes

just for the simple task of blurring an image. On a side note, I only managed to get the blurring outputting a greyscale image in the short period of time I spent using FIJI.

IVT was somewhat similar to FIJI in that it had a good range of high-level features, but was less impressive as a library. Despite following the example code I struggled to compile my own example and eventually gave up trying to get a working binary due to time constraints.

Having weighed up these libraries I it became fairly apparent that OpenCV would be the best choice, not only did it act as an easy to use library, it is also seems that is is one of the most prevalent libraries for Computer Vision. For the Kyffin Williams project OpenCV provides a lot of pre-built helpers which allow for very rapid production of the early elements of the project. For example it is able to handle loading images in different colour spaces, generating histograms of images and even comes with it's own machine learning libraries.

Having decided to go with OpenCV I was then faced with a choice of languages. OpenCV runs natively with C++ and has bindings for Python. Of these two languages I am slightly more familiar with C++ and I am also used to the syntax from programming Java continually for three years.

However, I wanted to take a rapid prototyping approach with this project; constantly adding new modules each week to build up a better and better system. C++ seemed too heavyweight for this approach, whilst Python seemed designed for it.

Another consideration I had was that the results of any analysis could take any form at all, from an array of numbers to a complex data structure returned by OpenCV. Trying to work around this with a statically typed language such as C++ could end up being difficult, especially when trying to keep a object-orientated approach. With a dynamically typed language such as Python it's a lot easier to pass these sorts of objects so long as the functions you pass to are expecting the right sort of object.

There may be some parts of the project which I might consider writing in C++ after prototyping in Python. An example of this would be if I wrote an algorithm which located strokes across the painting, as it may be used in other projects, not just my own.

2.2 Other Tools

Having used git for several personal projects I was keen to continue using it for my version control system. GitHub was a convenient hosting company to go with as they provide students with free private repositories and have very high uptime. It also means I am able to easily work on any system without difficulty but with the assurance of security. This is preferable to hosting my own repositories and running the risk of losing a lot of work.

Git is one of the nicer version control programs I have used and has some nice extensions (shell integration with zsh, git-flow for helping with the development process) and it's decentralised nature allows me to easily work where-ever I feel most comfortable.

```
import cv

im = cv.LoadImageM("tux.png")
blur = cv.CreateMat(im.rows, im.cols, im.type)
cv.Smooth(im, blur, cv.CV_BLUR, 9)
cv.SaveImage("tux_blurred_opencv.png", blur)
```

Figure 1: Using OpenCV to blur an image

```
package fiji;

import ij.IJ;
import ij.io.FileSaver;
import imagescience.feature.Smoother;
import imagescience.image.ColorImage;
import imagescience.image.Image;

public final class Blur {
    public static final String IMG_PATH = "tux.png";
    public static final String OUTPUT_PATH = "tux-fiji.png";

    public static void main(String[] args) {
        final Image image = new ColorImage(IJ.openImage(IMG_PATH));
        final Smoother s = new Smoother();
        final Image blurred = s.gauss(image, 9f);
        final FileSaver fs = new FileSaver(blurred.imageplus());
        fs.saveAsPng(OUTPUT_PATH);
    }
}
```

Figure 2: Using FIJI to blur an image

```
#include <stdio.h>
#include "Image/ByteImage.h"
#include "Image/ImageProcessor.h"

int main(int argc, char **argv) {
    if(argc != 2) {
        fprintf(stderr, "Usage: _ivt_input_output\n");
        return 1;
    }

    CByteImage in;
    CByteImage out;

    if(!in.LoadFromFile(argv[0])) {
        fprintf(stderr, "Unable_to_load_file_%s\n", argv[0]);
        return 2;
    }

    out = CByteImage(in);
    ImageProcessor::GaussianSmooth(&in, &out, 1.0f, 3);
    if(!out.SaveToFile(argv[1])) {
        fprintf(stderr, "Unable_to_save_to_file_%s\n", argv[1]);
    }
}
```

Figure 3: Using IVT to blur an image

Figure 4: Using VXT to blur an image

2.3 Design

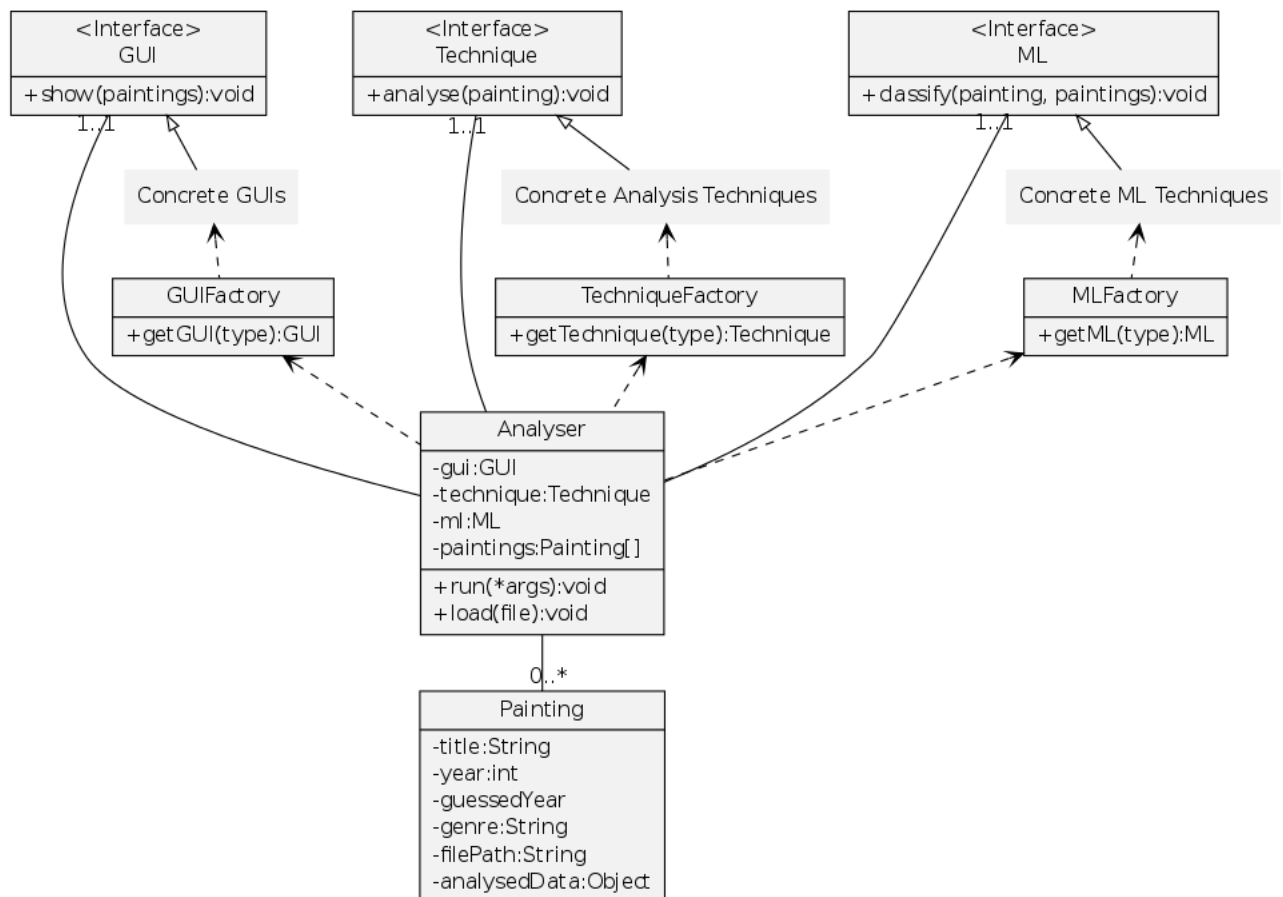


Figure 5: Initial Design

An initial UML class diagram is depicted by figure 5. This design will allow additional techniques, both analysis and machine learning based, to be added easily. Command line arguments, or later a GUI front-end, will be used to specify which techniques should be used.

The GUI elements depicted are used for visualising the analysed data or results of the machine learning algorithms. An example of this is a graphical representation of the colour space averages (see figure 7).

2.4 Colour Space Statistical Analysis

With the design solidified, I began to work on the statistical analysis of digital images. The easiest of these techniques is to read the image pixel by pixel, taking the various intensities at that pixel, then averaging them out over the whole image.

The most common colour space used is RGB (Red, Green, Blue), where each pixel holds three different intensities, one for each colour with a value of 0 to 255.

Using this analysis I plotted a graph of all paintings with a known year against the different intensities (see figure 6). As can be seen from this graph, the RGB values tend to fluctuate without much correlation between year.

Another popular colour space used in image processing is HSV (Hue, Saturation, Value) as it shows variations in colour, saturation and brightness separately (unlike RGB where all three values are affected by changes in brightness). It was a simple matter of changing the existing code for RGB so that it would analyse HSV values instead.

I also decided it would be an improvement to average all values in the same year together to help show changes as time goes on (see figure 7).

2.5 Histogram Analysis

With colour-space analysis complete, the next sensible step was to start generating colour histograms from Kyffin Williams paintings.

A histogram is a nice way of displaying the distribution of colour in a given image and is therefore a more useful representation of a painting than colour-space averages.

As with before, histograms can be taken from different colour spaces, currently only the program can only generate RGB histograms, but it is trivial to include HSV histograms too.

2.6 Nearest Neighbour Classification

Having finished the analysis techniques, then next part was to create a simple classification algorithm. The simplest of which is 1-Nearest Neighbour, that is; take the “Nearest” training example to the example to be classified and classify the year of this example with the year of the nearest example (see figure 8).

This simple algorithm can be expanded to be a K-Nearest Neighbour algorithm without too much effort (see figure 9). Instead of storing a single nearest neighbour you can simply store an list of them (making 1-Nearest Neighbour still work without issue) and then have some small amount of processing at the end to get a hold of the best neighbour. This processing will likely take the form of a simple statistic method such as the mean or the modal year.

2.7 Distance Measures

Working out the “Nearest” example to a given painting requires some way of measuring distance between the two.

One very simple way of doing this is Manhattan distance (figure 10). This distance measure can be improved further to euclidean distance (figure 11) without too much effort. There are numerous other methods of measuring the distance of points in space, including those which focus on the distance between colour histograms.

2.8 Meetings with Gareth Lloyd Roderick

At current I have had one meeting with both Lloyd and Hannah to discuss the current progress of the project and where we see the project progressing to. Lloyd was impressed by the current state of the project and was looking forward to see where it was leading to.

We also discussed the potential of getting a paper published out of this project.

2.9 Leave-one-out Cross Validation

The idea of leave-one-out cross validation is to take the entire training set, remove a single example and then to classify this example against the rest of the training set. You can then do this for every example in the training set and be able to plot a graph of the actual year against the classified year, aiming for the line $x = y$, an example of this is showing in figure 12. With this information we can also work out the correlation between the two and have a decent measure of how each analysis technique performs.

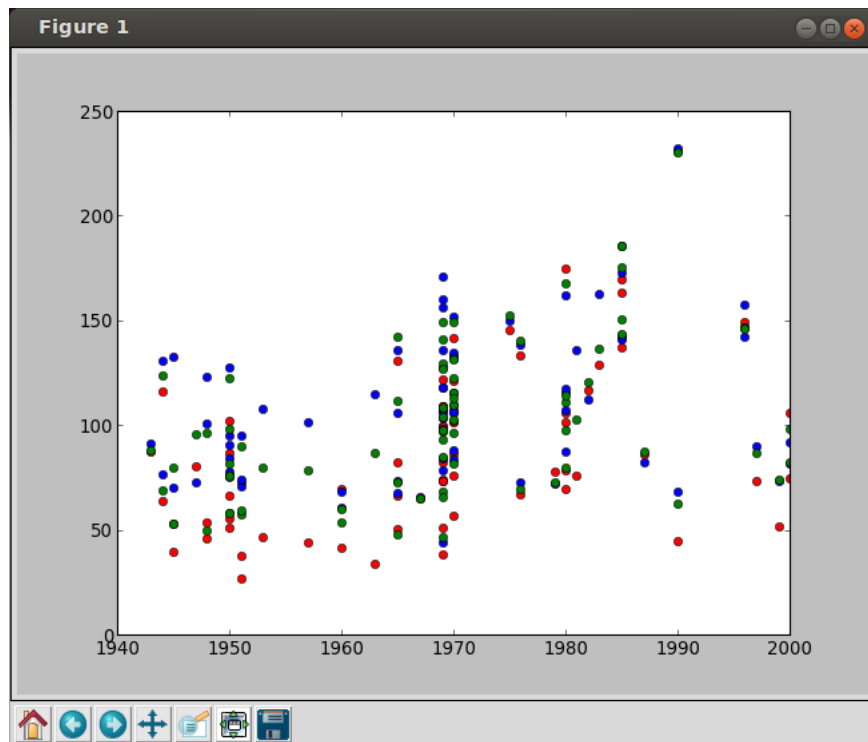


Figure 6: Mean RGB intensities by year.




Figure 7: Mean HSV intensities by year.

```

bestDistance  $\leftarrow \infty$ 
for  $x = 1 \rightarrow X$  do
  if  $\text{distance}(a, x) < \text{bestDistance}$  then
    bestDistance  $\leftarrow \text{distance}(a, x)$ 
  end if
end for

```

Where:

a is the example to classify.

X is the list of all training examples.

Figure 8: 1-Nearest Neighbour

```

kBest  $\leftarrow []$ 
for  $x = 1 \rightarrow X$  do
  cur  $\leftarrow \text{distance}(a, x)$ 
  for  $k = 1 \rightarrow K$  do
    if  $\text{cur} < \text{kBest}[k]$  then
      temp  $\leftarrow \text{kBest}[k]$ 
      kBest[ $k$ ]  $\leftarrow \text{cur}$ 
      cur  $\leftarrow \text{temp}$ 
    end if
  end for
end for

```

Where:

a is the example to classify.

X is the list of all training examples. K is the number of nearest neighbours to check.

Figure 9: K-Nearest Neighbour

2.10 Research into Stroke Analysis

Stroke analysis is one of the main goals for this project. It is quite apparent from looking at Kyffin Williams' paintings that his brushstrokes change over time, his early work having lots of smaller strokes over the canvas to large bold strokes in his later work.

The first paper I found relating to the analysis of brushstrokes involved moving a circular filter across the whole painting to find the ridges of strokes, then to fill any unbroken areas. They then shrunk these areas to a single pixel line and fitted a n^{th} order polynomial to this line [1]. This method seems fairly simplistic, but could be an interesting first step, but as it is more focused on authenticating paintings it may be of limited use.

Another method for stroke analysis has been published in the IEEE Transactions on Pattern Analysis and Machine Learning journal. This method is far more complex, but is able to extract and label individual brushstrokes. An interesting part of their findings was the ability to date some of Van Gogh's paintings to a known period in his career [2].

2.11 Technical Issues

Having engineered software before I know that the biggest issue with developing a piece of software is working with external APIs. With this project I can currently get away with only having to work with a few APIs; OpenCV being the main candidate. I have had some struggles trying to understand exactly what the API wants as inputs and returns as outputs, but that is more due to a slight lack in the area of computer vision and image processing rather than the API being complex.

$$d = \sum_{x=0}^X |a_x - b_x|$$

X : All dimensions present in both a and b .

a : The first point.

b : The second point.

Figure 10: Manhattan Distance

$$d = \sqrt{\sum_{x=0}^X (a_x - b_x)^2}$$

X : All dimensions present in both a and b .

a : The first point.

b : The second point.

Figure 11: Euclidean Distance

I am currently reading up on image processing and computer vision, current I'm only using free, online resources [3] and knowledge within the department.

OpenCV seems like a good library to work with at current, though I have yet to use it for any specialised image processing (i.e. I have only used built-in methods OpenCV provides by default). I currently don't see any issue as the python bindings for OpenCV use numpy in the background. Both of these libraries are widely used and open source. If I do happen to have any problems with either of these libraries there is a substantial community behind them and the option to fix any problems and give them back to the projects.

Another technical issue associated with this project is the resolution of the paintings in the current database. Most of the images are fairly small as they are intended to be used on a website rather than in an image processing application. There is potential that I may be able to get a hold of larger images for certain paintings, but if the techniques could be made to work on the current size of image it would probably be better for the project as a whole.

A problem associated with this is memory limits; for smaller images only a small amount of data will need to be stored, but for more complex techniques, large matrices may be generated and memory might become a precious resource. At current all analysis is run before classification but it may be necessary to only hold the analysis data for the current example to classify and the training example for which we are currently measuring the distance to. This would not be a difficult change to make, but may slow the program down as a whole.

3 Planning

3.1 Methodology

The methodology I am taking as part of this project is a mix of iterative development with rapid prototyping. This approach works well for this project as it is built up of a lot of small, isolated modules which perform discrete tasks.

Other models of development would likely slow the development of each of these modules down and are therefore less appropriate for this project. Some agile methodologies may also be a good approach, but as most agile methodologies are designed for teams rather than individual projects they are inappropriate.





Figure 12: Leave-one-out cross validation of actual year against classified year using RGB statistical analysis and 1-nearest neighbour classification.

3.2 Gantt Chart



3.3 Mid-term Demonstration

Annotated Bibliography

- [1] I. E. Berezhnuy, E. O. Postma, and H. J. van den Herik, "Authentic: Computerized brushstroke analysis," in *Multimedia and Expo, 2005. ICME 2005. IEEE International Conference on*. IEEE, July 2005, pp. 1586–1588. [Online]. Available: <http://dx.doi.org/10.1109/ICME.2005.1521739>

Defines a method of analysing brushstrokes by applying a circular filter across a digital image to pick up the ridges of a brushstroke. This can then be used to pick out individual brushstrokes in order to be able to fit a nth order polynomial to them. Though this paper focuses on authenticating Van Gogh's paintings, it could easily be applied to the work of Kyffin Williams and may allow for some interesting analysis.

- [2] J. Li, L. Yao, E. Hendriks, and J. Z. Wang, "Rhythmic brushstrokes distinguish van gogh from his contemporaries: Findings via automated brushstroke extraction," *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, vol. 34, no. 6, pp. 1159–1176, June 2012. [Online]. Available: <http://dx.doi.org/10.1109/TPAMI.2011.203>

Defines a complex method for analysing individual brush strokes which has been used to classify the period of two paintings by Van Gogh. This technique could be very powerful when applied to Kyffin Williams' work. This could be one of the most important techniques for the whole of the Kyffin Williams project.

- [3] S. J. D. Prince, *Computer vision : models, learning, and inference*. Cambridge University Press, 2012. [Online]. Available: <http://www.computervisionmodels.com/9781107011793>.

Learning reference for Computer Vision and Machine Learning.

- [4] D. Stork and M. Duarte, "Computer vision, image analysis, and master art: Part 3," *IEEE Multimedia*, vol. 14, no. 1, pp. 14–18, 2007. [Online]. Available: <http://dx.doi.org/10.1109/MMUL.2007.6>

Defines a potential method of analysing regions in a painting, this could be interesting to apply to Kyffin Williams' work as the regions in his earlier work are a lot less well defined as those in his later work.

- [5] D. G. Stork, "Computer vision and computer graphics analysis of paintings and drawings: An introduction to the literature computer analysis of images and patterns," ser. Lecture Notes in Computer Science, X. Jiang and N. Petkov, Eds. Berlin, Heidelberg: Springer Berlin / Heidelberg, 2009, vol. 5702, ch. 2, pp. 9–24. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-03767-2_2

Notes a lot of useful literature to look at and some useful terminology and ideas for analysis techniques too.